

# IN3050/IN4050 Mandatory Assignment 3: Unsupervised Learning

**Name:** Kevin Alexander Aslesen

**Username:** kevinaas

## Rules

Before you begin the exercise, review the rules at this website:

<https://www.uio.no/english/studies/examinations/compulsory-activities/mn-ifi-mandatory.html>, in particular the paragraph on cooperation. This is an individual assignment. You are not allowed to deliver together or copy/share source-code/answers with others. Read also the "Routines for handling suspicion of cheating and attempted cheating at the University of Oslo"

<https://www.uio.no/english/about/regulations/studies/studies-examinations/routines-cheating.html> By submitting this assignment, you confirm that you are familiar with the rules and the consequences of breaking them.

## Delivery

**Deadline:** Friday, April 26, 2024, 23:59

Your submission should be delivered in Devilry. You may redeliver in Devilry before the deadline, but include all files in the last delivery, as only the last delivery will be read. You are recommended to upload preliminary versions hours (or days) before the final deadline.

## What to deliver?

You are recommended to solve the exercise in a Jupyter notebook, but you might solve it in a Python program if you prefer.

If you choose Jupyter, you should deliver the notebook. You should answer all questions and explain what you are doing in Markdown. Still, the code should be properly commented. The notebook should contain results of your runs. In addition, you should make a pdf of your solution which shows the results of the runs.

If you prefer not to use notebooks, you should deliver the code, your run results, and a pdf-report where you answer all the questions and explain your work.

Your report/notebook should contain your name and username.

Deliver one single zipped folder (.zip, .tgz or .tar.gz) which contains your complete solution.

Important: if you weren't able to finish the assignment, use the PDF report/Markdown to elaborate on what you've tried and what problems you encountered. Students who have made an effort and attempted all parts of the assignment will get a second chance even if they fail initially. This exercise will be graded PASS/FAIL.

## Goals of the exercise

This exercise has three parts. The first part is focused on Principal Component Analysis (PCA). You will go through some basic theory, and implement PCA from scratch to do compression and visualization of data.

The second part focuses on clustering using K-means. You will use `scikit-learn` to run K-means clustering, and use PCA to visualize the results.

The last part ties supervised and unsupervised learning together in an effort to evaluate the output of K-means using a logistic regression for multi-class classification approach.

The master students will also have to do one extra part about tuning PCA to balance compression with information lost.

## Tools

You may freely use code from the weekly exercises and the published solutions. In the first part about PCA you may **NOT** use ML libraries like `scikit-learn`. In the K-means part and beyond we encourage the use of `scikit-learn` to iterate quickly on the problems.

# Principal Component Analysis (PCA)

In this section, you will work with the PCA algorithm in order to understand its definition and explore its uses. Some sources for more information on PCA are:

- The syllabus book by Marsland has an overview of the mathematics and coding involved on page 136-137.
- For a more intuitive explanation of PCA, there are many good explanations online, like [this one](#).
- If you are puzzled by what the covariance matrix is, and how it relates to PCA, [this video](#) may be useful.

## Implementation: how is PCA implemented?

Here we implement the basic steps of PCA and we assemble them.

### Importing libraries

We start importing the `numpy` library for performing matrix computations, the `pyplot` library for plotting data, and the `syntheticdata` module to import synthetic data.

```
In [ ]: import numpy as np
import matplotlib.pyplot as plt

import syntheticdata
```

## Centering the Data

Implement a function with the following signature to center the data. Remember that every *feature* should be centered.

```
In [ ]: def center_data(A):
    # INPUT:
    # A      [NxM] numpy data matrix (N samples, M features)
    #
    # OUTPUT:
    # X      [NxM] numpy centered data matrix (N samples, M features)

    mean = np.mean(A, axis=0)
    X = A - mean

    return X
```

Test your function checking the following assertion on *testcase*:

```
In [ ]: testcase = np.array([[3., 11., 4.3], [4., 5., 4.3], [5., 17., 4.5], [4, 13., 4.4]
answer = np.array([[-1., -0.5, -0.075], [0., -6.5, -0.075], [1., 5.5, 0.125], [0., 0., 0.]])
np.testing.assert_array_almost_equal(center_data(testcase), answer)
```

## Computing Covariance Matrix

Implement a function with the following signature to compute the covariance matrix. In order to get this at the correct scale, divide by  $N - 1$ , not  $N$ . Do not use `np.cov()`.

```
In [ ]: def compute_covariance_matrix(A):
    # INPUT:
    # A      [NxM] numpy data matrix (N samples, M features)
    #
    # OUTPUT:
    # C      [MxM] numpy covariance matrix (M features, M features)

    N = A.shape[0]
    X = center_data(A)
    C = (X.T @ X) / (N-1)

    return C
```

Test your function checking the following assertion on *testcase*:

```
In [ ]: test_array = np.array([[22., 11., 5.5], [10., 5., 2.5], [34., 17., 8.5], [28., 14., 7.5]])
answer = np.cov(np.transpose(test_array))
to_test = compute_covariance_matrix(test_array)
```

```
np.testing.assert_array_almost_equal(to_test, answer)
```

## Computing eigenvalues and eigenvectors

Use the linear algebra package of `numpy` and its function `np.linalg.eig()` to compute eigenvalues and eigenvectors. Notice that we take the real part of the eigenvectors and eigenvalues. The covariance matrix *should* be a symmetric matrix, but the actual implementation in `compute_covariance_matrix()` can lead to small round off errors that lead to tiny imaginary additions to the eigenvalues and eigenvectors. These are purely numerical artifacts that we can safely remove.

**Note:** If you decide to NOT use `np.linalg.eig()` you must make sure that the eigenvalues you compute are of unit lenght!

```
In [ ]: def compute_eigenvalue_eigenvectors(A):
    # INPUT:
    # A      [DxD] numpy matrix
    #
    # OUTPUT:
    # eigval    [D] numpy vector of eigenvalues
    # eigvec   [DxD] numpy array of eigenvectors

    eigval, eigvec = np.linalg.eig(A)

    # Numerical roundoff can lead to (tiny) imaginary parts. We correct that here
    eigval = eigval.real
    eigvec = eigvec.real

    return eigval, eigvec
```

Test your function checking the following assertion on `testcase`:

```
In [ ]: testcase = np.array([[2, 0, 0], [0, 5, 0], [0, 0, 3]])
answer1 = np.array([2., 5., 3.])
answer2 = np.array([[1., 0., 0.], [0., 1., 0.], [0., 0., 1.]])
x, y = compute_eigenvalue_eigenvectors(testcase)
np.testing.assert_array_almost_equal(x, answer1)
np.testing.assert_array_almost_equal(y, answer2)
```

## Sorting eigenvalues and eigenvectors

Implement a function with the following signature to sort eigenvalues and eigenvectors.

Remember that eigenvalue `eigval[i]` corresponds to eigenvector `eigvec[:, i]`.

```
In [ ]: def sort_eigenvalue_eigenvectors(eigval, eigvec):
    # INPUT:
    # eigval    [D] numpy vector of eigenvalues
    # eigvec   [DxD] numpy array of eigenvectors
    #
    # OUTPUT:
    # sorted_eigval    [D] numpy vector of eigenvalues
    # sorted_eigvec   [DxD] numpy array of eigenvectors
```

```

# Get indices
indices = np.argsort(eigval) # smallest to biggest eigenvalue
indices = indices[::-1]      # biggest to smallest eigenvalue

# Sort with indices
sorted_eigval = eigval[indices]
sorted_eigvec = eigvec[:, indices]

return sorted_eigval, sorted_eigvec

```

Test your function checking the following assertion on *testcase*:

```
In [ ]: testcase = np.array([[2, 0, 0], [0, 5, 0], [0, 0, 3]])
answer1 = np.array([5., 3., 2.])
answer2 = np.array([[0., 0., 1.], [1., 0., 0.], [0., 1., 0.]])
x, y = compute_eigenvalue_eigenvectors(testcase)
x, y = sort_eigenvalue_eigenvectors(x, y)
np.testing.assert_array_almost_equal(x, answer1)
np.testing.assert_array_almost_equal(y, answer2)
```

## PCA Algorithm

Implement a function with the following signature to compute PCA using the functions implemented above.

```
In [ ]: def pca(A, m):
    # INPUT:
    # A      [NxM] numpy data matrix (N samples, M features)
    # m      integer number denoting the number of learned features (m <= M)
    #
    # OUTPUT:
    # pca_eigvec   [Mxm] numpy matrix containing the eigenvectors (M dimensions)
    # P            [Nxm] numpy PCA data matrix (N samples, m features)

    # Center data and compute covariance matrix
    X = center_data(A)
    C = compute_covariance_matrix(X)

    # Get eigenvalues and eigenvectors and sort from biggest to smallest
    eigval, eigvec = compute_eigenvalue_eigenvectors(C)
    sorted_eigval, sorted_eigvec = sort_eigenvalue_eigenvectors(eigval, eigvec)

    pca_eigvec = sorted_eigvec[:, :m]

    P = (pca_eigvec.T @ X.T).T

    return pca_eigvec, P
```

Test your function checking the following assertion on *testcase*:

```
In [ ]: import pickle
testcase = np.array([[22., 11., 5.5], [10., 5., 2.5], [34., 17., 8.5]])
x, y = pca(testcase, 2)

answer1_file = open('PCAanswer1.pkl', 'rb')
answer2_file = open('PCAanswer2.pkl', 'rb')
```

```
answer1 = pickle.load(answer1_file)
answer2 = pickle.load(answer2_file)

test_arr_x = np.sum(np.abs(np.abs(x) - np.abs(answer1)), axis=0)
np.testing.assert_array_almost_equal(test_arr_x, np.zeros(2))

test_arr_y = np.sum(np.abs(np.abs(y) - np.abs(answer2)))
np.testing.assert_almost_equal(test_arr_y, 0)
```

## Understanding: how does PCA work?

We now use the PCA algorithm you implemented on a toy data set in order to understand its inner workings.

### Loading the data

The module *syntheticdata* provides a small synthetic dataset of dimension [100x2] (100 samples, 2 features).

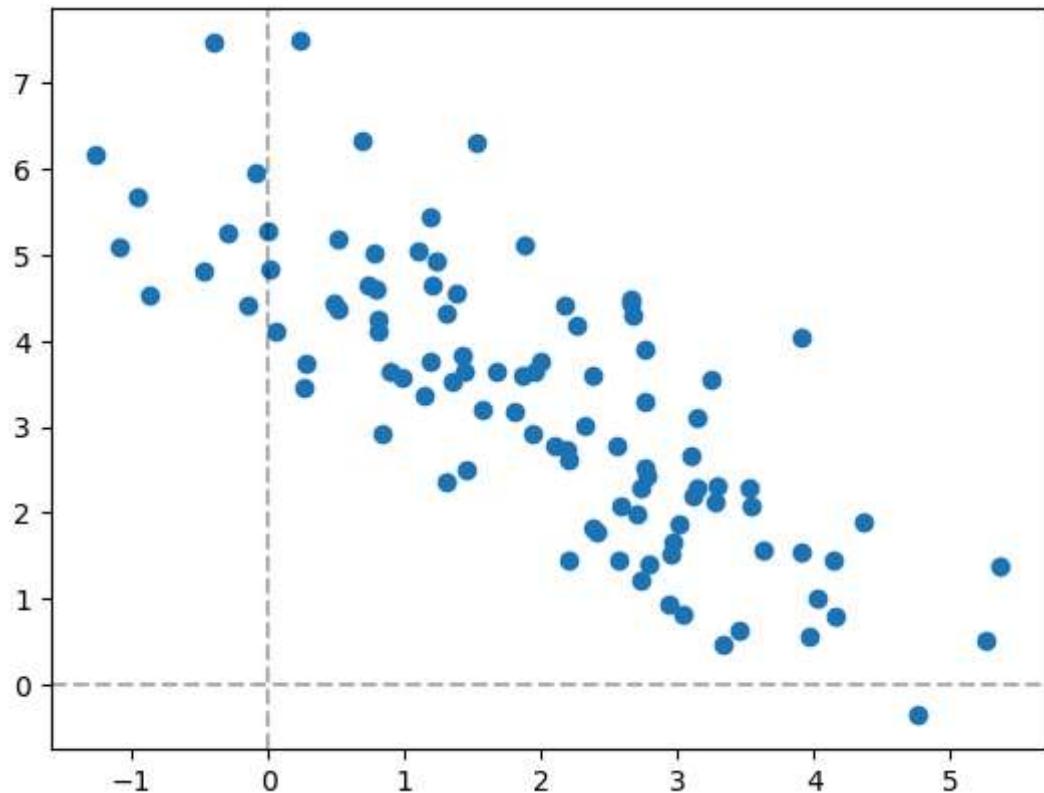
```
In [ ]: X = syntheticdata.get_synthetic_data1()
```

### Visualizing the data

Visualize the synthetic data using the function *scatter()* from the *matplotlib* library.

```
In [ ]: plt.scatter(X[:, 0], X[:, 1])
plt.axhline(0, color = 'k', linestyle = '--', alpha=0.3)
plt.axvline(0, color = 'k', linestyle = '--', alpha=0.3)
```

```
Out[ ]: <matplotlib.lines.Line2D at 0x237e9cff9d0>
```

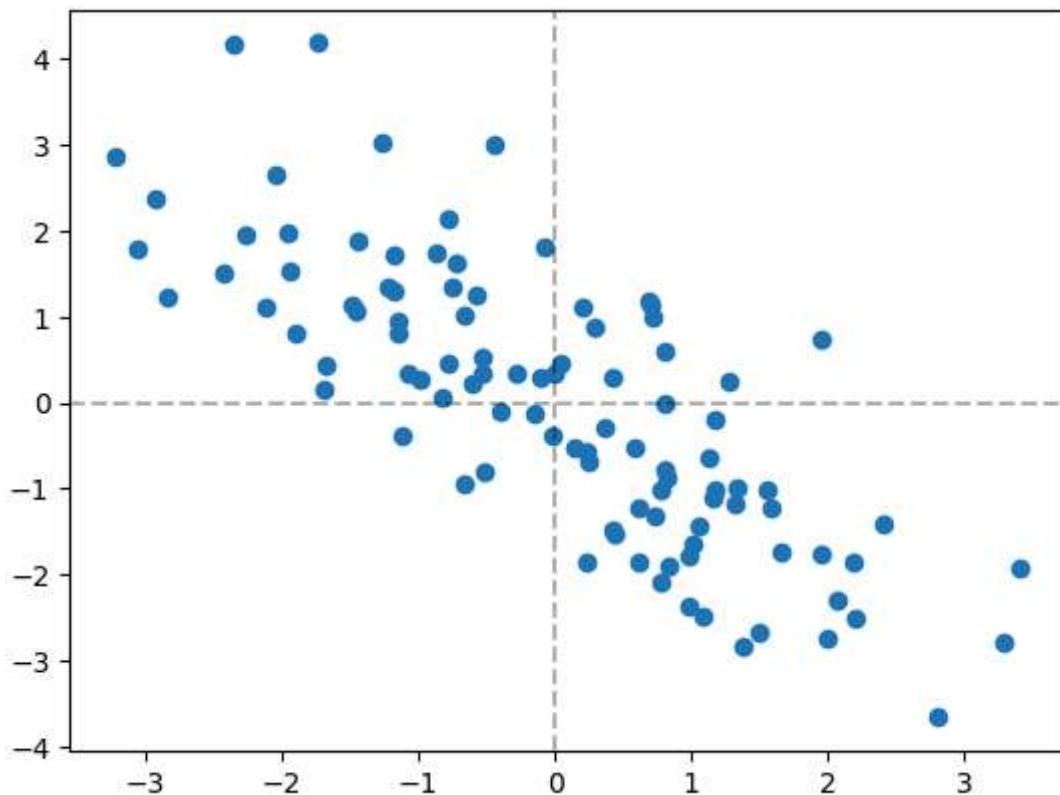


## Visualize the centered data

Notice that the data visualized above is not centered on the origin (0,0). Use the function defined above to center the data, and the replot it.

```
In [ ]: X = center_data(X)
plt.scatter(X[:, 0], X[:, 1])
plt.axhline(0, color = 'k', linestyle = '--', alpha=0.3)
plt.axvline(0, color = 'k', linestyle = '--', alpha=0.3)
```

```
Out[ ]: <matplotlib.lines.Line2D at 0x237ebd7cb50>
```



## Visualize the first eigenvector

Visualize the vector defined by the first eigenvector. To do this you need:

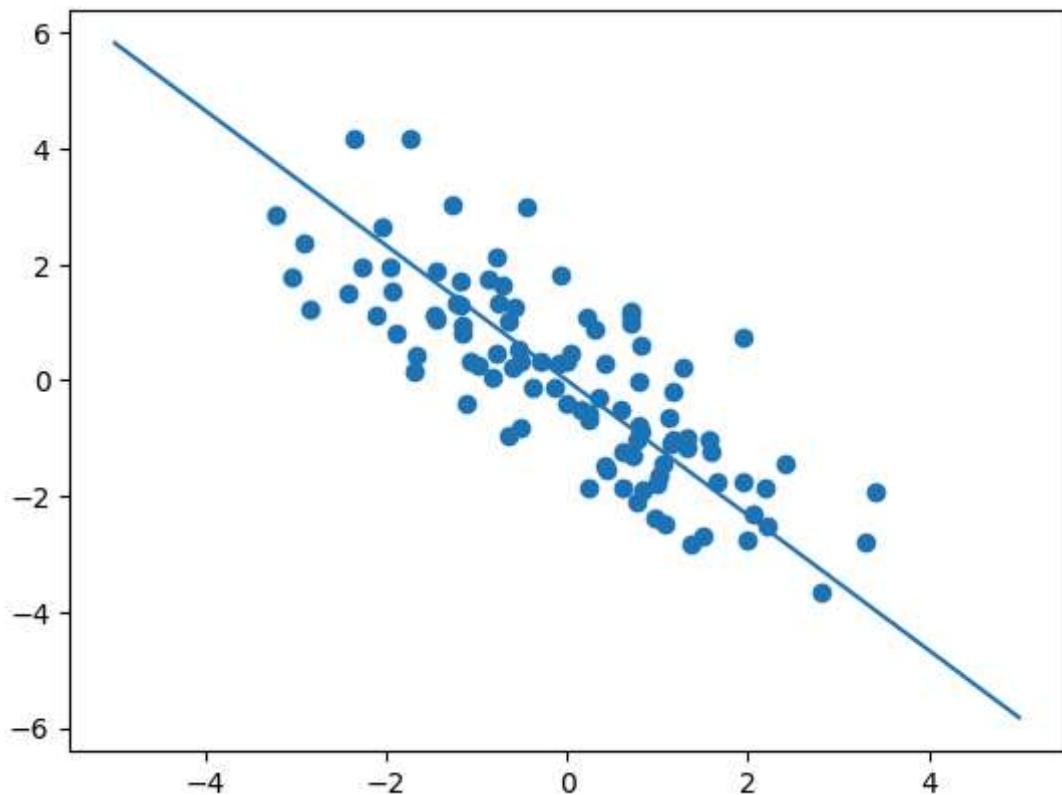
- Use the *PCA()* function to recover the eigenvectors
- Plot the centered data as done above
- The first eigenvector is a 2D vector  $(x_0, y_0)$ . This defines a vector with origin in  $(0,0)$  and head in  $(x_0, y_0)$ . Use the function *plot()* from matplotlib to plot a line over the first eigenvector.

```
In [ ]: pca_eigvec, _ = pca(X, 2)
first_eigvec = pca_eigvec[0]

plt.scatter(X[:, 0], X[:, 1])

x = np.linspace(-5, 5, 1000)
y = first_eigvec[1] / first_eigvec[0] * x
plt.plot(x, y)
```

```
Out[ ]: [<matplotlib.lines.Line2D at 0x237ec021250>]
```

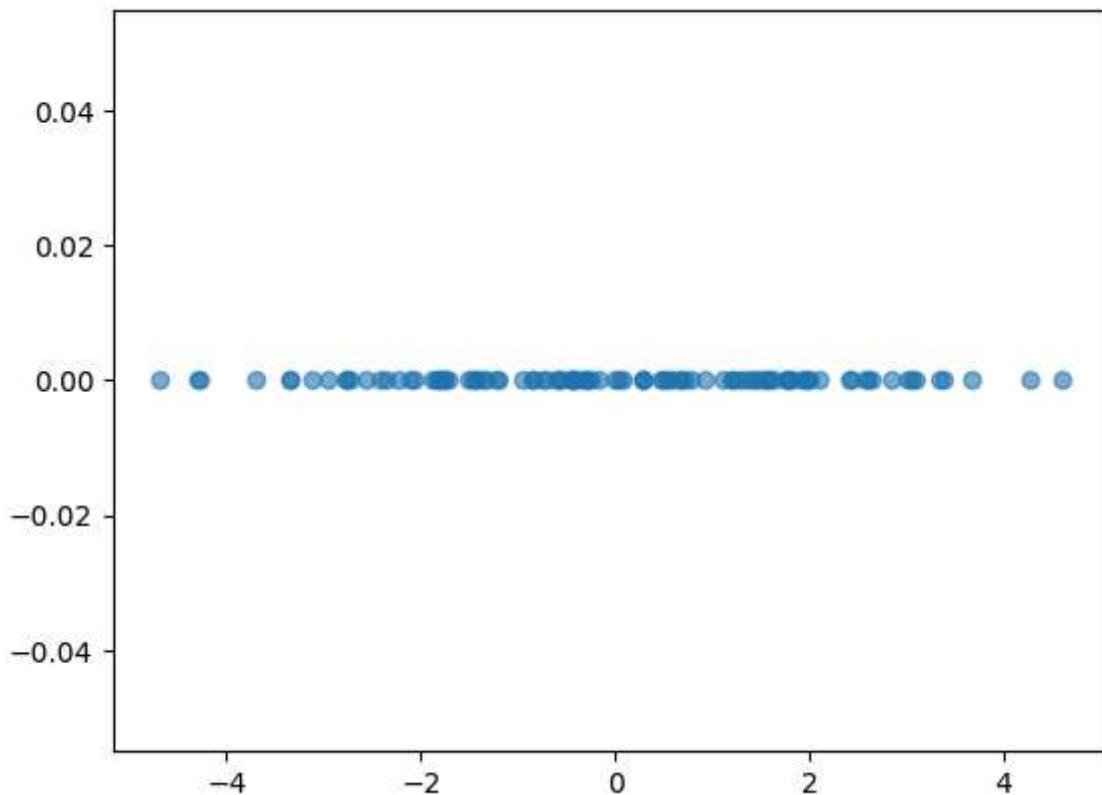


## Visualize the PCA projection

Finally, use the `PCA()` algorithm to project on a single dimension and visualize the result using again the `scatter()` function.

```
In [ ]: __, P = pca(X, 1)
plt.scatter(P[:, 0], np.zeros(len(P[:, 0])), alpha=0.6)
```

```
Out[ ]: <matplotlib.collections.PathCollection at 0x237ec068750>
```



## Evaluation: when are the results of PCA sensible?

So far we have used PCA on synthetic data. Let us now imagine we are using PCA as a pre-processing step before a classification task. This is a common setup with high-dimensional data. We explore when the use of PCA is sensible.

### Loading the first set of labels

The function `get_synthetic_data_with_labels1()` from the module `syntheticdata` provides a first labeled dataset.

```
In [ ]: X, y = syntheticdata.get_synthetic_data_with_labels1()
```

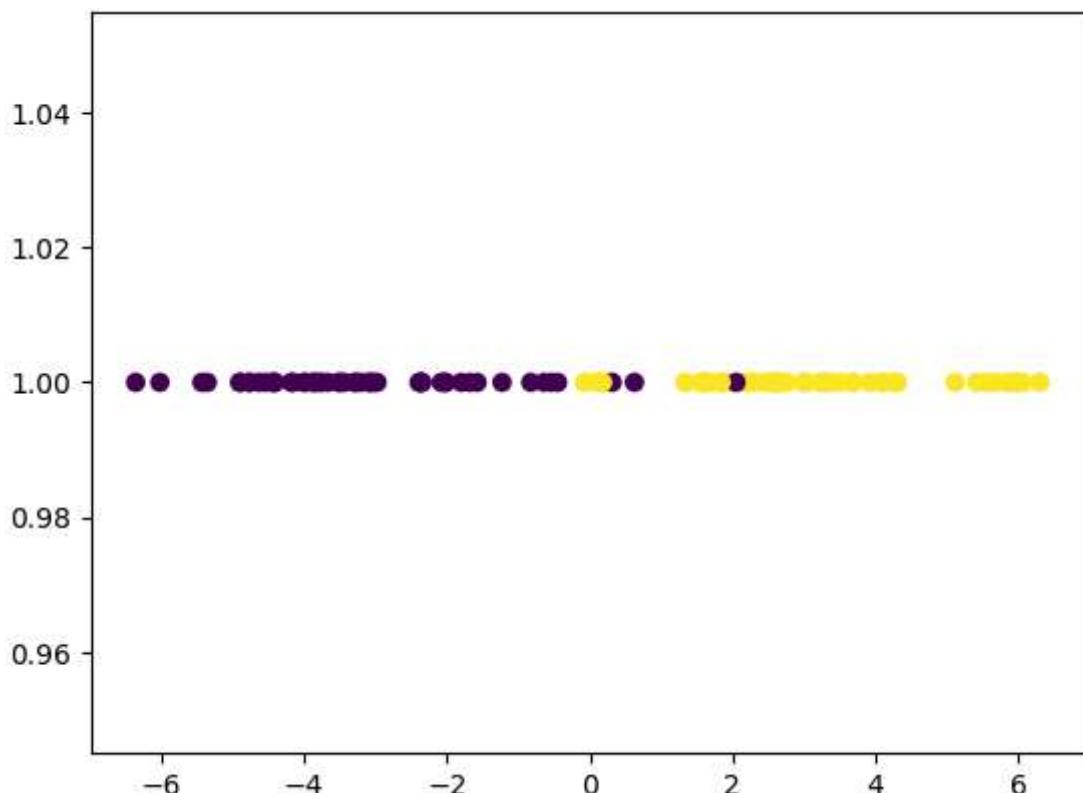
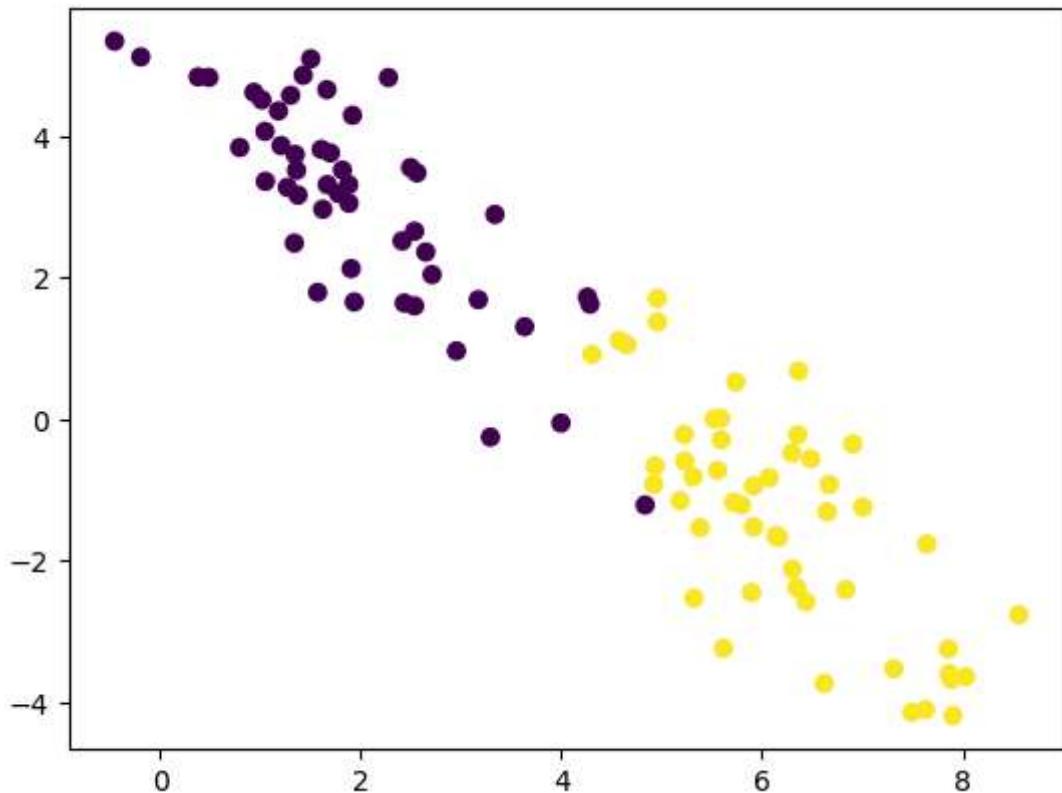
### Running PCA

Process the data using the PCA algorithm and project it in one dimension. Plot the labeled data using `scatter()` before and after running PCA. Comment on the results.

```
In [ ]: plt.scatter(X[:, 0], X[:, 1], c=y[:, 0])

plt.figure()
_, P = pca(X, 1)
plt.scatter(P[:, 0], np.ones(P.shape[0]), c=y[:, 0])
```

```
Out[ ]: <matplotlib.collections.PathCollection at 0x237ec14c1d0>
```



**Comment:** The two classes (purple and yellow) are easy to separate in both the two and one dimensional scatter plots.

## Loading the second set of labels

The function `get_synthetic_data_with_labels2()` from the module `syntheticdata` provides a second labeled dataset.

```
In [ ]: X, y = syntheticdata.get_synthetic_data_with_labels2()
```

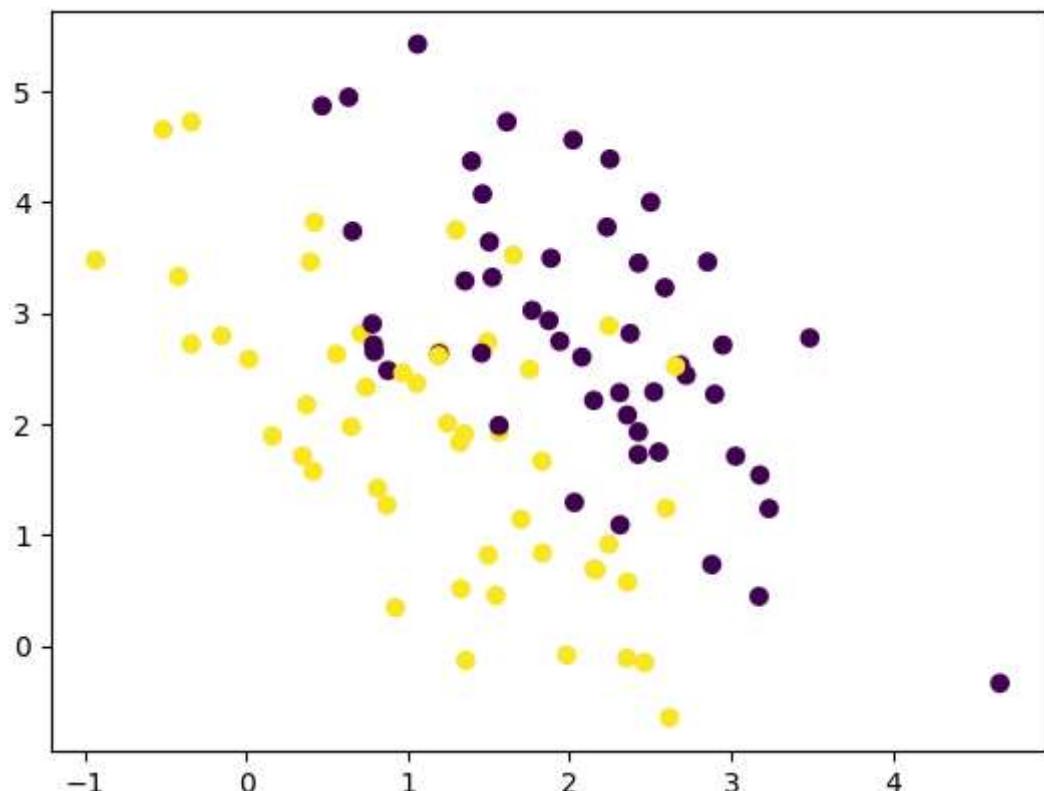
## Running PCA

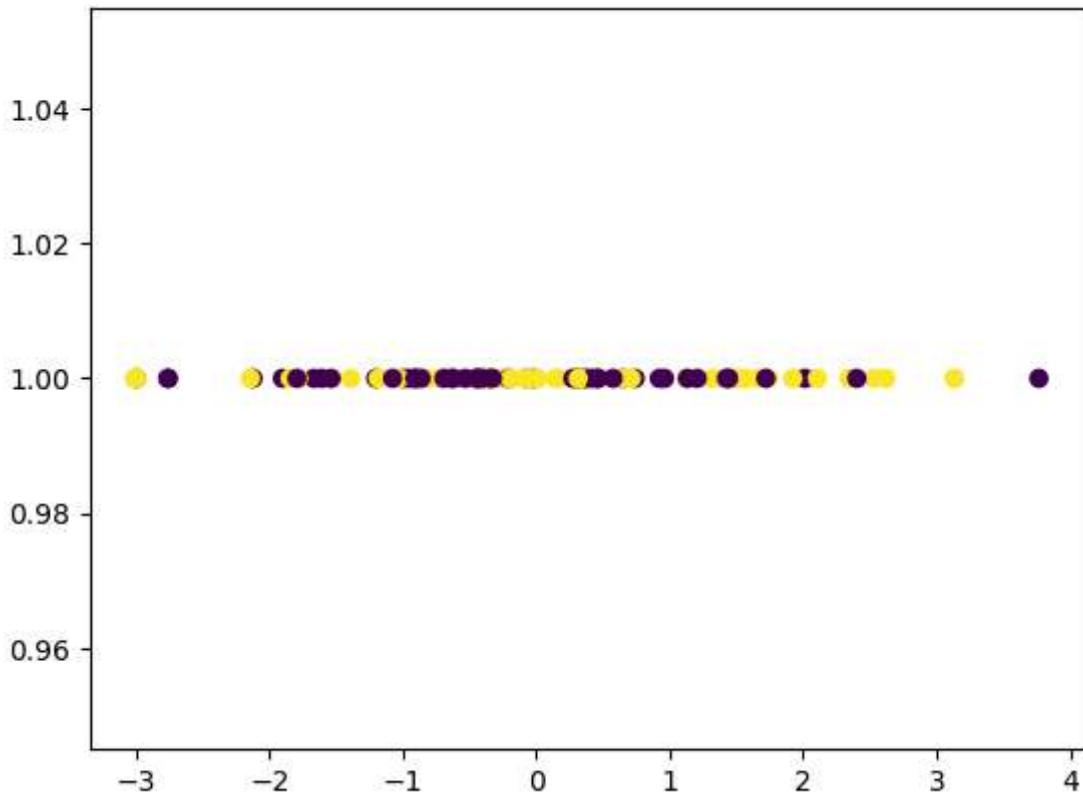
As before, process the data using the PCA algorithm and project it in one dimension. Plot the labeled data using `scatter()` before and after running PCA. Comment on the results.

```
In [ ]: plt.scatter(X[:, 0], X[:, 1], c=y[:, 0])

plt.figure()
_, P = pca(X, 2)
plt.scatter(P[:, 0], np.ones(P.shape[0]), c=y[:, 0])
```

```
Out[ ]: <matplotlib.collections.PathCollection at 0x237ec4374d0>
```





**Comment:** In this case, the two classes are easy to separate with a diagonal line in the two dimensional plot, while they are much harder to separate in the one dimensional plot.

How would the result change if you were to consider only the second eigenvector? What about if you were to consider both eigenvectors?

**Answer:** The second eigenvector would point in the orthogonal direction, meaning where the data is less variable. The two classes seem to be more separated in this direction, so the resulting one dimensional plot would be more separable than the one dimensional plot above. If we were to consider both eigenvectors, then we would end up with a two dimensional plot that is more separable than the one dimensional plot above as well.

## Case study 1: PCA for visualization

We now consider the *iris* dataset, a simple collection of data ( $N=150$ ) describing iris flowers with four ( $M=4$ ) features. The features are: Sepal Length, Sepal Width, Petal Length and Petal Width. Each sample has a label, identifying each flower as one of 3 possible types of iris: Setosa, Versicolour, and Virginica.

Visualizing a 4-dimensional dataset is impossible; therefore we will use PCA to project our data in 2 dimensions and visualize it.

### Loading the data

The function `get_iris_data()` from the module `syntheticdata` returns the *iris* dataset. It returns a data matrix of dimension [150x4] and a label vector of dimension [150].

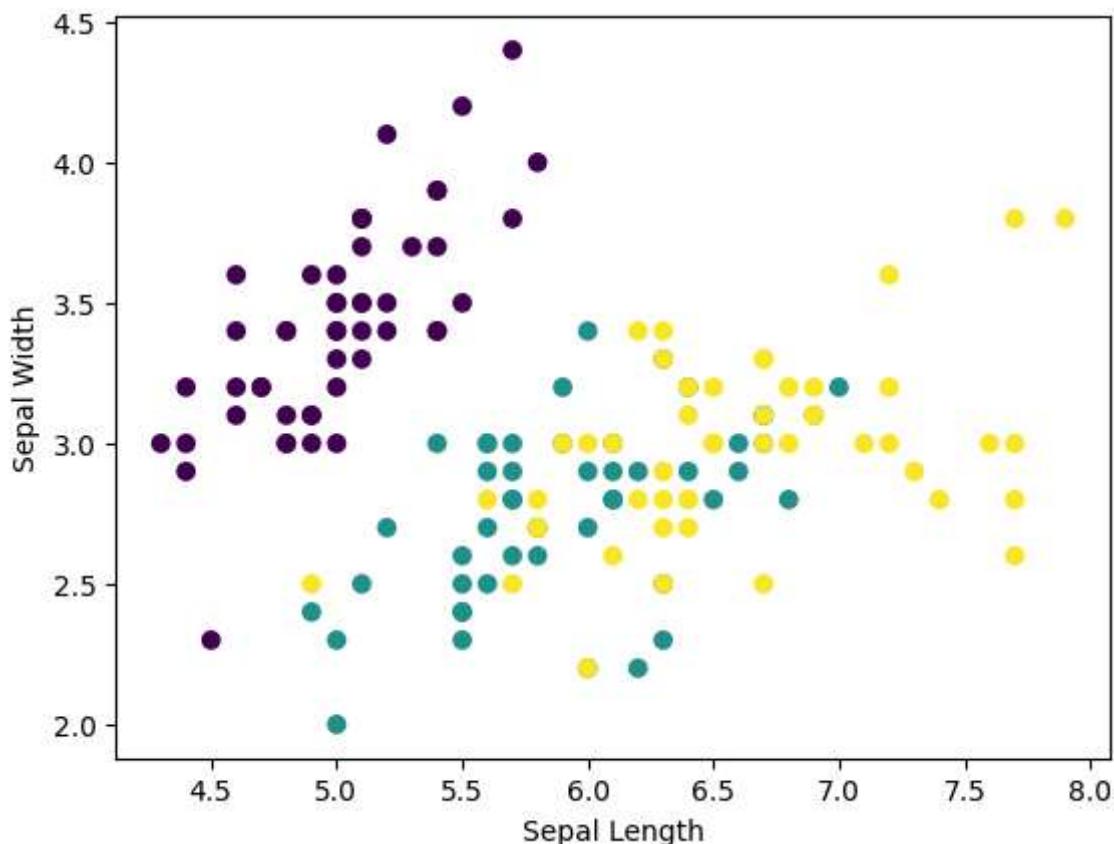
```
In [ ]: X, y = syntheticdata.get_iris_data()
```

## Visualizing the data by selecting features

Try to visualize the data (using label information) by randomly selecting two out of the four features of the data. You may try different pairs of features.

```
In [ ]: # Get the features
sepal_length = X[:, 0]
sepal_width = X[:, 1]
petal_length = X[:, 2]
petal_width = X[:, 3]

# Plot
plt.scatter(sepal_length, sepal_width, c=y)
plt.xlabel('Sepal Length')
plt.ylabel('Sepal Width')
plt.show()
```

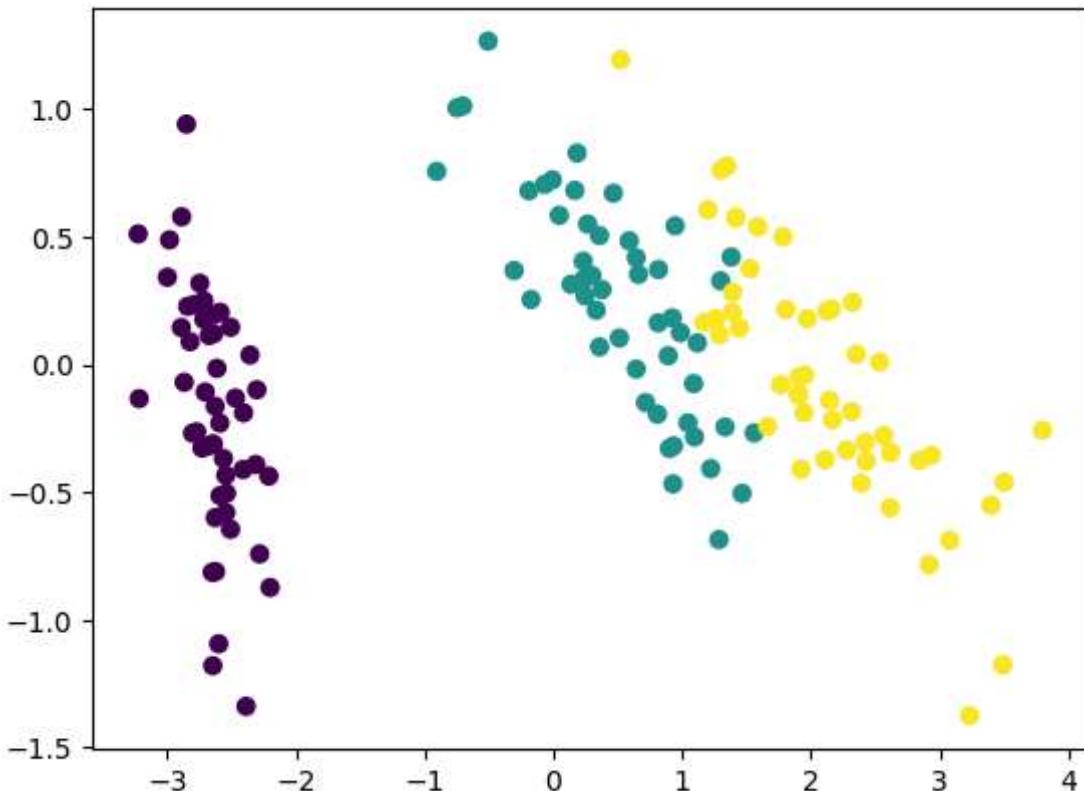


## Visualizing the data by PCA

Process the data using PCA and visualize it (using label information). Compare with the previous visualization and comment on the results.

```
In [ ]: _, P = pca(X, 2)
plt.scatter(P[:, 0], P[:, 1], c=y)
```

```
plt.show()
```



**Comment:** This is much more separable than the first plot, especially for the green and yellow classes compared to before PCA.

## Case study 2: PCA for compression

We now consider the *faces in the wild (lfw)* dataset, a collection of pictures ( $N=1280$ ) of people. Each pixel in the image is a feature ( $M=2914$ ).

### Loading the data

The function `get_lfw_data()` from the module `syntheticdata` returns the `lfw` dataset. It returns a data matrix of dimension [1280x2914] and a label vector of dimension [1280]. It also returns two parameters,  $h$  and  $w$ , reporting the height and the width of the images (these parameters are necessary to plot the data samples as images). Beware, it might take some time to download the data. Be patient :)

```
In [ ]: x, y, h, w = syntheticdata.get_lfw_data()
```

### Inspecting the data

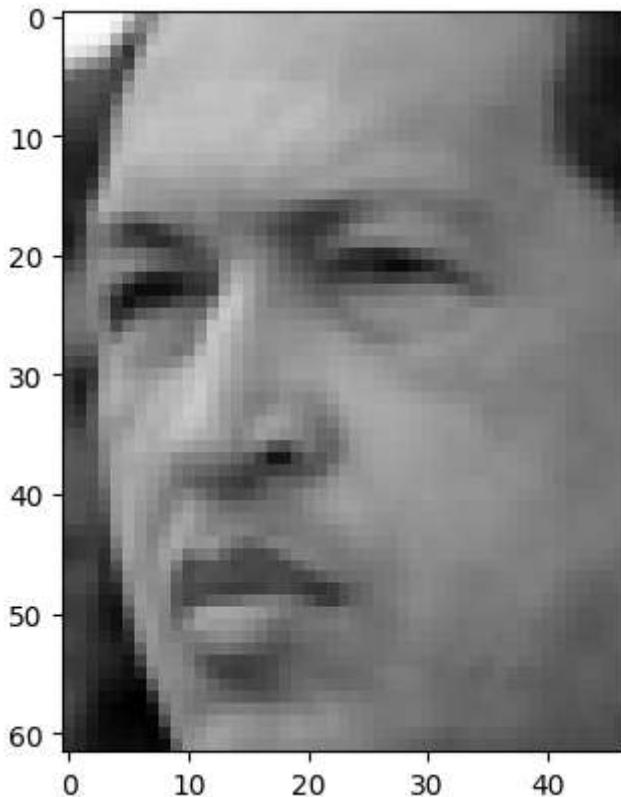
Choose one datapoint to visualize (first coordinate of the matrix  $X$ ) and use the function `imshow()` to plot and inspect some of the pictures.

Notice that `imshow` receives as a first argument an image to be plot; the image must be provided as a rectangular matrix, therefore we reshape a sample from the matrix  $X$  to

have height  $h$  and width  $w$ . The parameter  $cmap$  specifies the color coding; in our case we will visualize the image in black-and-white with different gradations of grey.

```
In [ ]: plt.imshow(X[0, :].reshape((h, w)), cmap=plt.cm.gray)
```

```
Out[ ]: <matplotlib.image.AxesImage at 0x237ec5619d0>
```



## Implementing a compression-decompression function

Implement a function that first uses PCA to project samples in low-dimensions, and then reconstruct the original image.

*Hint:* Most of the code is the same as the previous PCA() function you implemented.

```
In [ ]: def encode_decode_pca(A, m):
    # INPUT:
    # A      [NxM] numpy data matrix (N samples, M features)
    # m      integer number denoting the number of learned features (m <= M)
    #
    # OUTPUT:
    # Ahat [NxM] numpy PCA reconstructed data matrix (N samples, M features)

    # Center data and compute covariance matrix
    X = center_data(A)
    C = compute_covariance_matrix(X)

    # Get eigenvalues and eigenvectors and sort from biggest to smallest
    eigval, eigvec = compute_eigenvalue_eigenvectors(C)
    sorted_eigval, sorted_eigvec = sort_eigenvalue_eigenvectors(eigval, eigvec)

    pca_eigvec = sorted_eigvec[:, :m]
```

```
P = (pca_eigvec.T @ X.T).T
Ahat = P @ pca_eigvec.T

return Ahat
```

## Compressing and decompressing the data

Use the implemented function to encode and decode the data by projecting on a lower dimensional space of dimension 200 ( $m=200$ ).

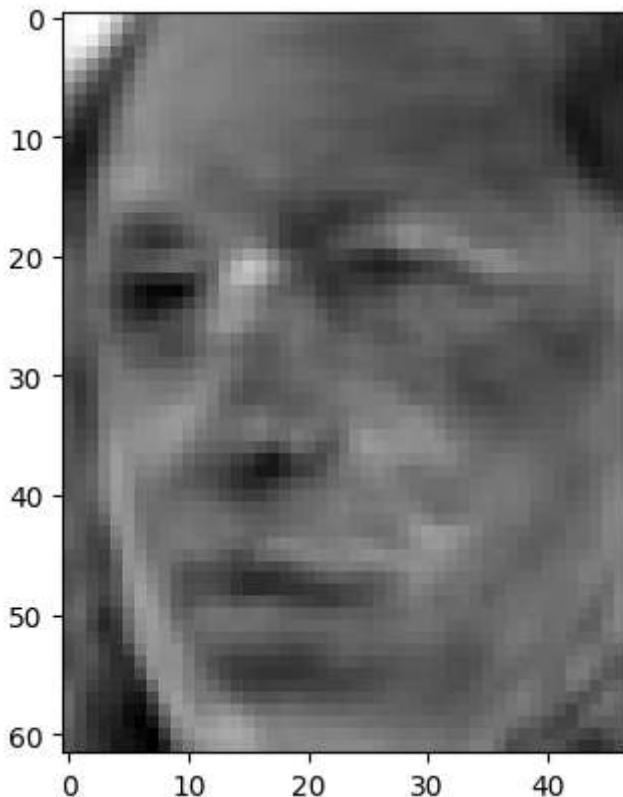
```
In [ ]: Xhat = encode_decode_pca(X, 200)
```

## Inspecting the reconstructed data

Use the function `imshow` to plot and compare original and reconstructed pictures. Comment on the results.

```
In [ ]: plt.imshow(Xhat[0, :].reshape((h, w)), cmap=plt.cm.gray)
```

```
Out[ ]: <matplotlib.image.AxesImage at 0x237ec5ceb90>
```



**Comment:** The reconstructed picture have clearly lost some of the details that were present in the original picture, but it is still possible to identify that there is a man in the picture. The reconstructed picture also looks much darker than the original.

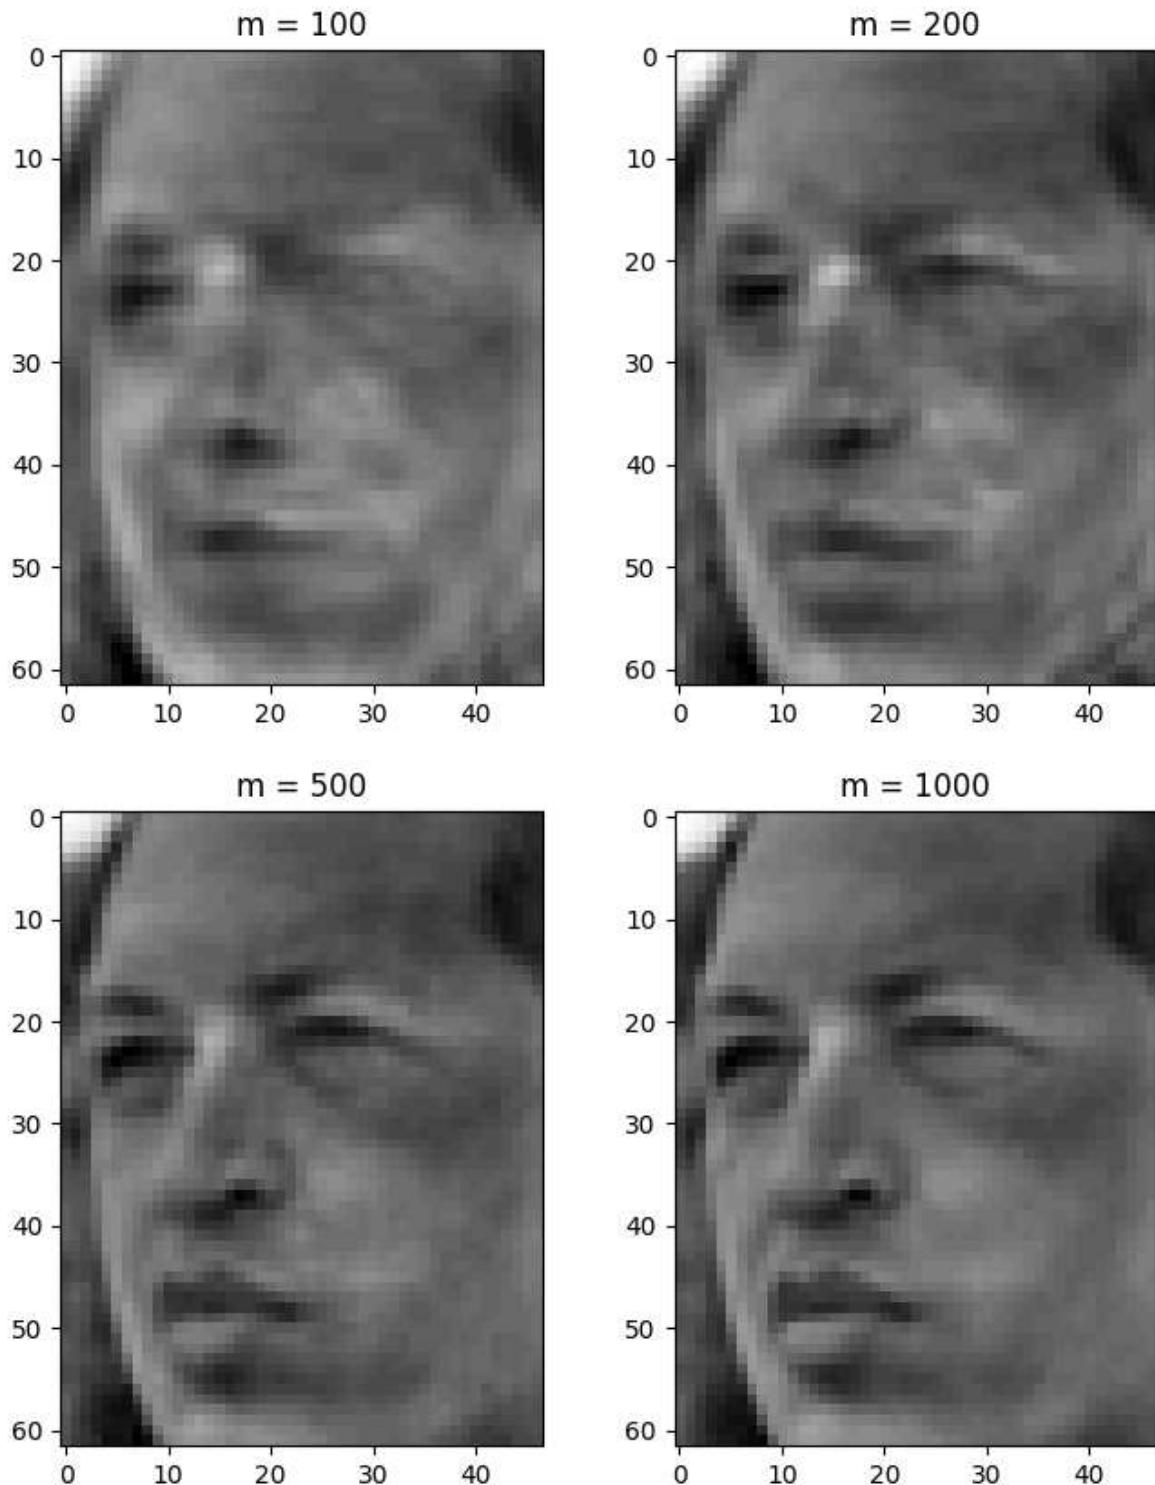
## Evaluating different compressions

Use the previous setup to generate compressed images using different values of low dimensions in the PCA algorithm (e.g.: 100, 200, 500, 1000). Plot and comment on the

results. Try to use `plt.subplot(n_rows, n_cols, position)`, in addition to titles, to get a nice plot.

```
In [ ]: plt.figure(figsize=(8, 10))

for i,m in enumerate([100, 200, 500, 1000]):
    plt.subplot(2, 2, i+1)
    X_hat = encode_decode_pca(X, m)
    plt.imshow(X_hat[0, :].reshape((h, w)), cmap=plt.cm.gray)
    plt.title(f'm = {m}'')
```



**Comment:** The details in the picture compressed to 100 dimensions are quite bad. One could mistakenly think the person is much older than he actually is or that he might

suffer from a skin disease. As the dimensions become bigger, the details become more prominent. The picture with 1000 dimensions looks almost identical to the original picture, maybe it even looks better because of the higher contrast in the picture.

## Master Students: PCA Tuning

If we use PCA for compression or decompression, it may be not trivial to decide how many dimensions to keep. In this section we review a principled way to decide how many dimensions to keep.

The number of dimensions to keep is the only *hyper-parameter* of PCA. A method designed to decide how many dimensions/eigenvectors is the *proportion of variance*:

$$\text{POV} = \frac{\sum_{i=1}^m \lambda_i}{\sum_{j=1}^M \lambda_j},$$

where  $\lambda$  are eigenvalues,  $M$  is the dimensionality of the original data, and  $m$  is the chosen lower dimensionality.

Using the *POV* formula we may select a number  $M$  of dimensions/eigenvalues so that the proportion of variance is, for instance, equal to 95%.

Implement a new PCA for encoding and decoding that receives in input not the number of dimensions for projection, but the amount of proportion of variance to be preserved.

```
In [ ]: def encode_decode_pca_with_pov(A, p):
    # INPUT:
    # A      [NxM] numpy data matrix (N samples, M features)
    # p      float number between 0 and 1 denoting the POV to be preserved
    #
    # OUTPUT:
    # Ahat   [NxM] numpy PCA reconstructed data matrix (N samples, M features)
    # m      integer reporting the number of dimensions selected

    # Center data and compute covariance matrix
    X = center_data(A)
    C = compute_covariance_matrix(X)

    # Get eigenvalues and eigenvectors and sort from biggest to smallest
    eigval, eigvec = compute_eigenvalue_eigenvectors(C)
    sorted_eigval, sorted_eigvec = sort_eigenvalue_eigenvectors(eigval, eigvec)

    # Start from M and subtract with 1 for each POV evaluation that are bigger than
    m = X.shape[1]  # M
    POV = 1
    while (POV > p):
        POV = np.sum(sorted_eigval[:m]) / np.sum(sorted_eigval)
        m -= 1

    pca_eigvec = sorted_eigvec[:, :m]

    P = (pca_eigvec.T @ X.T).T
    Ahat = P @ pca_eigvec.T
```

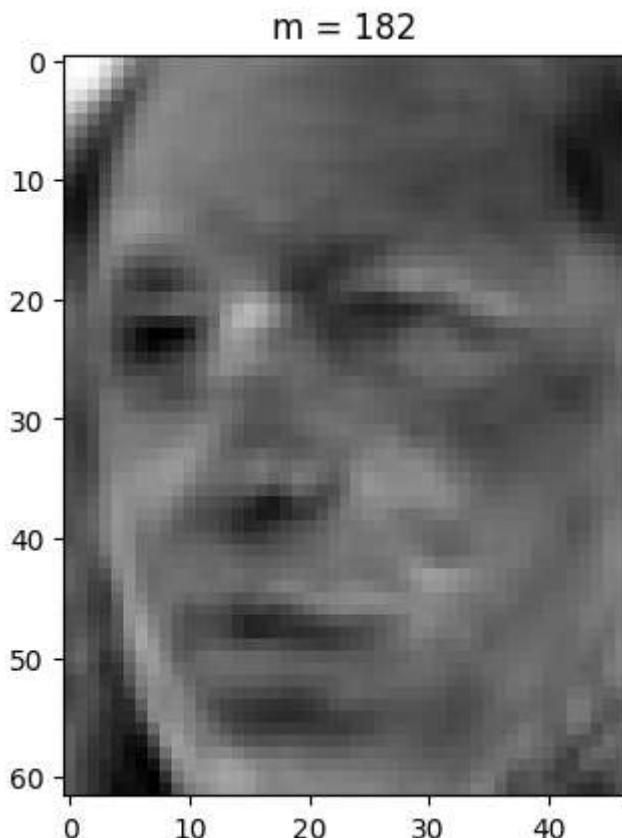
```
    return Ahat, m
```

Import the `lfw` dataset using the `get_lfw_data()` in `syntheticdata`. Use the implemented function to encode and decode the data by projecting on a lower dimensional space such that `POV=0.95`. Use the function `imshow` to plot and compare original and reconstructed pictures. Comment on the results.

```
In [ ]: X, y, h, w = syntheticdata.get_lfw_data()
```

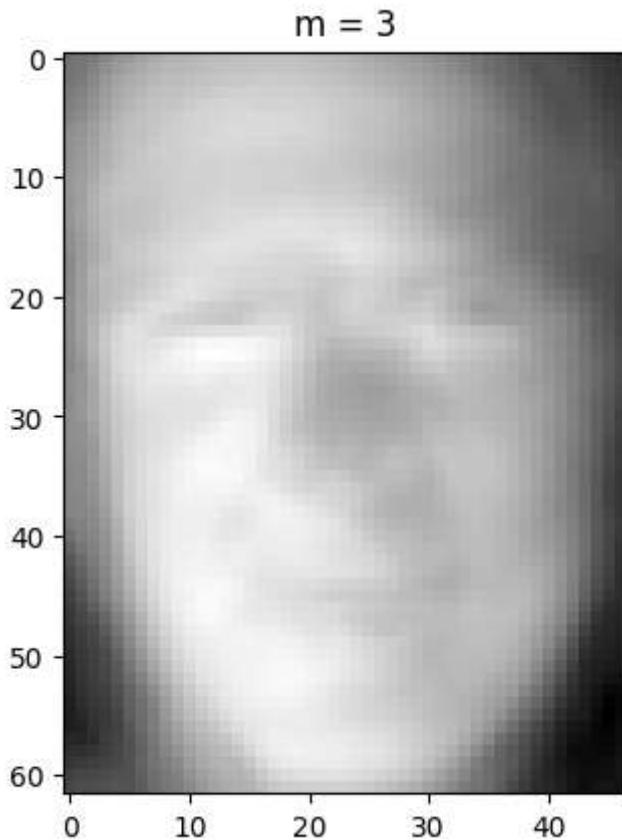
```
In [ ]: Xhat, m = encode_decode_pca_with_pov(X, 0.95)
```

```
In [ ]: # Plot the images here
plt.imshow(Xhat[0, :].reshape((h, w)), cmap=plt.cm.gray)
plt.title(f'm = {m}')
plt.show()
```



**Comment:** This picture show fewer details than the original, undoubtly, but I imagine I would still be able to indentify the person in the picture. We see that it only needed 182 dimensions to accomplish this, which is impressive considering the original picture had almost 3000 dimensions (features).

```
In [ ]: Xhat, m = encode_decode_pca_with_pov(X, 0.5)
plt.imshow(Xhat[0, :].reshape((h, w)), cmap=plt.cm.gray)
plt.title(f'm = {m}')
plt.show()
```



This picture will haunt me the next couple of weeks.

## K-Means Clustering (Bachelor and master students)

In this section you will use the *k-means clustering* algorithm to perform unsupervised clustering. Then you will perform a qualitative assessment of the results.

### Importing scikit-learn library

We start importing the module `sklearn.cluster.KMeans` from the standard machine learning library `scikit-learn`.

```
In [ ]: from sklearn.cluster import KMeans
```

### Loading the data

We will use once again the *iris* data set. The function `get_iris_data()` from the module `syntheticdata` returns the *iris* dataset. It returns a data matrix of dimension [150x4] and a label vector of dimension [150].

```
In [ ]: X, y = syntheticdata.get_iris_data()
```

### Projecting the data using PCA

To allow for visualization, we project our data in two dimensions as we did previously. This step is not necessary, and we may want to try to use *k-means* later without the PCA pre-processing. However, we use PCA, as this will allow for an easy visualization.

```
In [ ]: _, P = pca(X, 2)
```

## Running k-means

We will now consider the *iris* data set as an unlabeled set, and perform clustering to this unlabeled set. We can compare the results of the clustering to the labeled classes.

Use the class *KMeans* to fit and predict the output of the *k-means* algorithm on the projected data. Run the algorithm using the following values of  $k = \{2, 3, 4, 5\}$ .

```
In [ ]: k_values = [2, 3, 4, 5]
y_hats = []

for i in range(len(k_values)):
    KM = KMeans(k_values[i], random_state=0, n_init="auto") # Fill in correct
    yhat = KM.fit_predict(P) # Store the prediction
    y_hats.append(yhat)
```

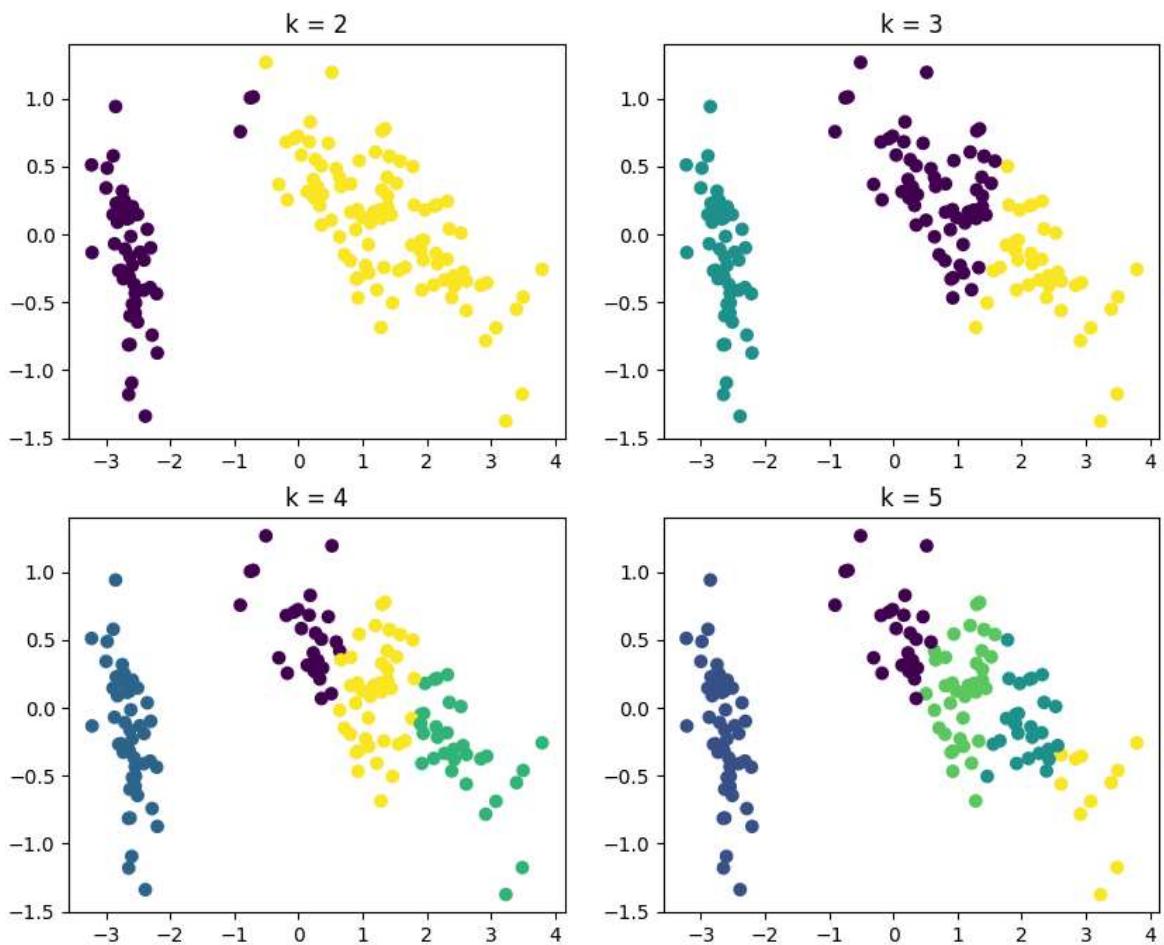
## Qualitative assessment

Plot the results of running the k-means algorithm, compare with the true labels, and comment. Try to use `plt.subplot(n_rows, n_cols, position)` with titles to make a nice plot.

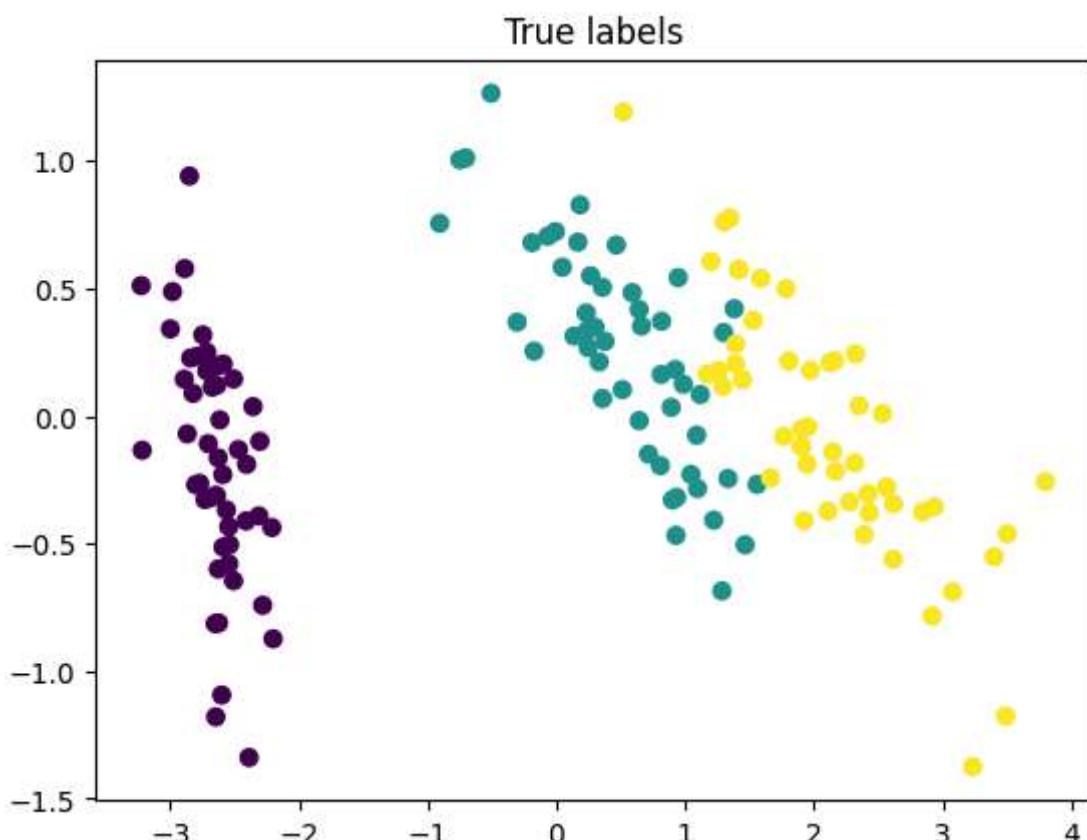
**Hint:** Plot the first and second dimension of `P` using `plt.scatter()`, and set the keyword argument `c` to the predictions made with *KMeans*.

```
In [ ]: plt.figure(figsize=(10, 8))

for i in range(len(k_values)):
    plt.subplot(2, 2, i+1)
    plt.scatter(P[:, 0], P[:, 1], c=y_hats[i])
    plt.title(f'k = {k_values[i]}')
```



```
In [ ]: _, P = pca(X, 2)
plt.scatter(P[:, 0], P[:, 1], c=y)
plt.title('True labels')
plt.show()
```



**Comment:** Comparing the clusters for different k-values with the true labels above, we see that k=3 gives the most similar result to the true labels. This is expected since the iris-dataset is split into three classes.

## Quantitative Assessment of K-Means (Bachelor and master students)

We used k-means for clustering and we assessed the results qualitatively by visualizing them. However, we often want to be able to measure in a quantitative way how good the clustering was. To do this, we will use a classification task to **evaluate numerically how good the learned clusters are for all the different values of k you used above (2 to 5).**

Informally, our evaluation will work as follows: For each of our clusterings (k=2 to k=5), we will try to learn a mapping from the identified clusters to the correct labels of our datapoints. The reason this can be a sensible evaluation, is that to learn a good mapping, we have to have identified clusters that correspond to the actual classes in our data. We will in other words train 4 different classification models (one for each k value), where the input to our classifier is the cluster each datapoint belongs to, and the target is the correct class for this datapoint. In other words, **we aim to learn to classify datapoints as well as possible with the only information available to the classifier being the cluster that datapoint belongs to.** For some values of k, we will get poorly performing classifiers, indicating that this clustering has not revealed the correct class division in our data.

In practice, here is what you will do: Reload the *iris* dataset. Import a standard `LogisticRegression` classifier from the module `sklearn.linear_model`. Use the k-means representations learned previously (`yhat2, ..., yhat5`) and the true label to train the classifier. Evaluate your model on the training data (we do not have a test set, so this procedure will assess the model fit instead of generalization) using the `accuracy_score()` function from the `sklearn.metrics` module. Plot a graph showing how the accuracy score varies when changing the value of k. Comment on the results.

- Train a Logistic regression model using the first two dimensions of the PCA of the iris data set as input, and the true classes as targets.
- Report the model fit/accuracy on the training set.
- For each value of K:
  - One-Hot-Encode the classes output by the K-means algorithm.
  - Train a Logistic regression model on the K-means classes as input vs the real classes as targets.
  - Calculate model fit/accuracy vs. value of K.
- Plot your results in a graph and comment on the K-means fit.

```
In [ ]: from sklearn.linear_model import LogisticRegression  
from sklearn import metrics
```

```

import pandas as pd

lr = LogisticRegression()
lr.fit(P, y)
acc = metrics.accuracy_score(lr.predict(P), y)
print(f'Accuracy on PCA vs true classes: {acc:.3f}')

accs = []
for i in range(len(k_values)):
    # One-hot encode the classes given by the kmean algorithm
    one_hot = pd.get_dummies(y_hats[i])

    # Train a Logistic regression model on the one-hot as input and the real cl
    lr_kmean = LogisticRegression()
    lr_kmean.fit(one_hot, y)

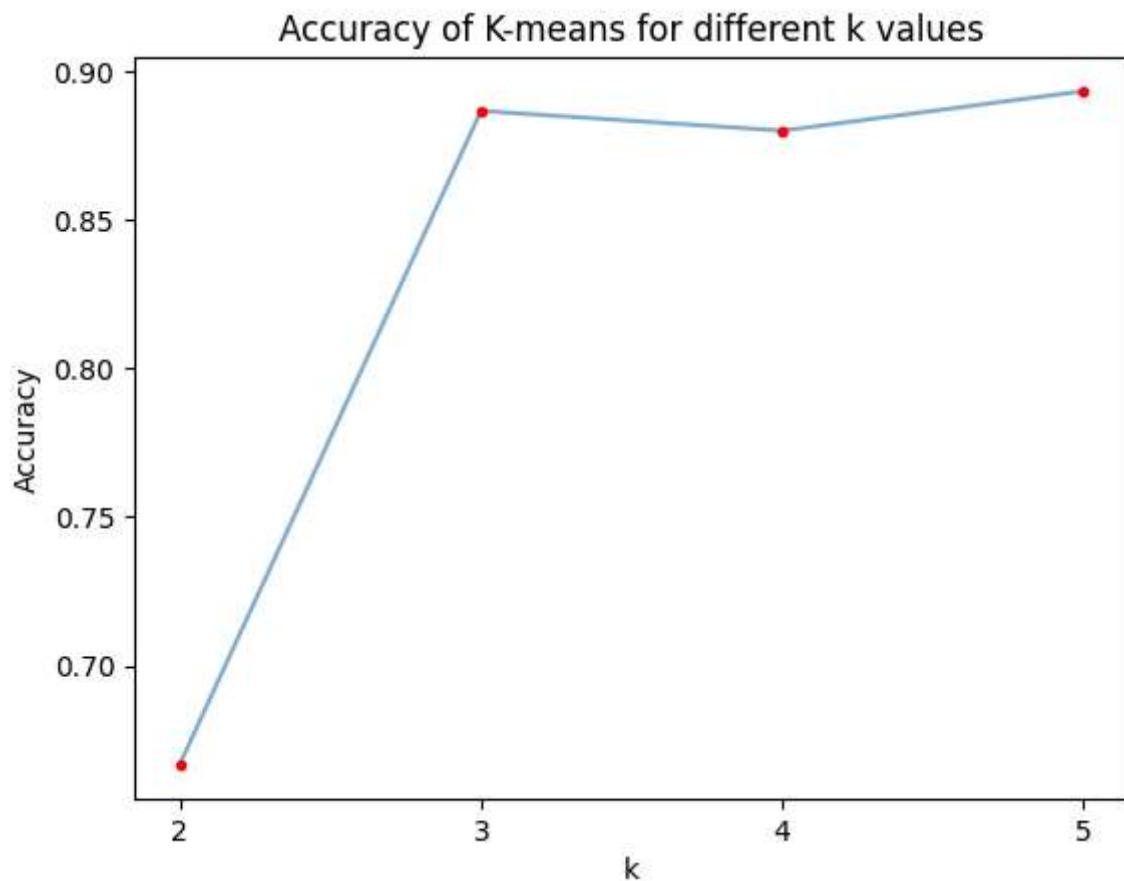
    # Calculate the accuracy and store in list above
    acc_mean = metrics.accuracy_score(lr_kmean.predict(one_hot), y)
    accs.append(acc_mean)

plt.plot(k_values, accs, alpha=0.6)
plt.plot(k_values, accs, 'r.')
plt.xlabel('k')
plt.ylabel('Accuracy')
plt.title('Accuracy of K-means for different k values')
plt.xticks(k_values)
plt.show()

print('-- K-means accuracies --')
for i in range(len(k_values)):
    print(f'k={k_values[i]}: {accs[i]:.3f}')

```

Accuracy on PCA vs true classes: 0.967



```
-- K-means accuracies --
k=2: 0.667
k=3: 0.887
k=4: 0.880
k=5: 0.893
```

**Comment:** Firstly, we see that the PCA algorithm with 2 dimensions predicts the true classes very well with an accuracy of 0.967, which I find impressive since the original data had 4 dimensions (features). The K-means algorithm however does not perform as well, especially in the case where k=2, in which the accuracy is an nonimpressive 0.667. It picks up however when we go up to k=3, giving an accuracy of 0.887. In this case there are as many clusters as there are true classes, for which I believed would score the highest accuracy score. It turns out I was wrong however, since we see that after the accuracy slightly drops when k=4, the accuracy jumps up to a staggering 0.907 when k=5. I guess this could make sense if the two separate classes in k=5 (yellow and dark blue) together fit the yellow points in the *true labels* plot better than the single class in k=3 (dark green). This illustrated the importance of a quantitative evaluation compared to an only qualitative evaluation.

## Conclusions

In this notebook we studied **unsupervised learning** considering two important and representative algorithms: **PCA** and **k-means**.

First, we implemented the PCA algorithm step by step; we then run the algorithm on synthetic data in order to see its working and evaluate when it may make sense to use it and when not. We then considered two typical uses of PCA: for **visualization** on the *iris* dataset, and for **compression-decompression** on the *lfw* dataset.

We then moved to consider the k-means algorithm. In this case we used the implementation provided by *scikit-learn* and we applied it to another prototypical unsupervised learning problem: **clustering**; we used *k-means* to process the *iris* dataset and we evaluated the results visually.

In the final part, we considered two additional questions that may arise when using the above algorithms. For PCA, we considered the problem of **selection of hyper-parameters**, that is, how we can select the hyper-parameter of our algorithm in a reasonable fashion. For k-means, we considered the problem of the **quantitative evaluation** of our results, that is, how can we measure the performance or usefulness of our algorithms.